


LOCALSY

So I decided to take some images since I don't know your hardware and I assume you wouldn't wanna run this locally yet. In case you do wanna download and run it, for lower spec hardware it'd probably download the Gemma E4B vision+text gguf model. When I did build this AI assistant, I was also building it for my friends who don't have that much hardware and I primarily developed and tested it with smaller models and made sure it worked with those first.

General Chat:


In the first image you can see how the AI reacts to a query that implies you want up to date information. The AI is told that if your question is something that must have time sensitive information, then it will do a search and so we see it done here. There are also a bunch of security features gating prompt injections, but even if you would get prompt injections this AI doesn't have access to a terminal and can't run commands on your computer. The most it can do is read from your computer but not write to it, and only within a limited upload folder scope and its own database.



Anyway, having that out of the way, on the first image you can see:

- Execution trace. This is for seeing what the AI model does in the background.
- Thinking process, streamed. If the task you gave the AI is taking a while or if the AI is generally big and it takes a while for it to do stuff, instead of being forced to stare at a spinner you can at least read the streamed thought process.

- On the top of the chat you can also see a “context data” and that’s one of the most important things for local AI imo.



So here we can see the finished response. Of course, the thought process is streamed and so is the response of the AI once it starts. No loading spinner and nothing appears out of thin air all of a sudden.

- You can see the context data having the f1 driver standing search query and it contains the raw token amount (this is an estimated number but more or less okay).
- Towards the bottom of the image you can see which sources we extracted data from and if you click them, it opens up the exact website.
- Then finally you can see some metrics. We can see that Qwen 3.8 27b is loaded (on the left bottom under system health) and in the chat bubble as the author of it. This gives you insight into how the model actually performs on the hardware. We can see how long it was thinking for, what tokens per second i got, what the time to first token was and how much of the conversation is cached in kv values (any tool usage injects external data into the model resulting in cache bust, but I built this system to bust the cache as little as possible. Busting the cache results in pre-fill time from scratch and a long waiting time so its important we keep cache intact as much as possible)
- On top you can see the contex window bar and the max context window the AI model was loaded with. There is also a condense chat button that reduces the bloat of the conversation and packs it into as few tokens as possible giving you continuity. You can also manually handle context and more on that on the next image.

search_web: US military pilot officer pay scale 2026, aviation incentive pay military pilots, military pilot com... x

SOURCES

Defense Finance and Accounting Service > MilitaryMembers > payentitlements > Pay Tables
www.dfas.mil

2026 Military Pay Chart 3.8% (All Pay Grades)
www.navycs.com

Defense Finance and Accounting Service > MilitaryMembers > payentitlements > Pay Tables > AVIP5
www.dfas.mil

Military Pilot Pay & Benefits: What Aviators Really Earn
bogidope.com

How Much Do Air Force Pilots Make in a Year? – USAMM
www.usamm.com

Pilot Salary USA 2026: All Major Airlines Compared (\$105K–\$419K)
airmappr.com

RAW EVIDENCE ~15872 tokens

2026 Military Pay Chart 3.8% (All Pay Grades)

2026 U.S. Military Basic Pay Charts

2026 U.S. Military Basic Pay Charts On December 18, 2025 , President Trump signed the National Defense Authorization Act for Fiscal Year 2026 (S. 5009) into law. The 2026 military pay raise of 3.8 percent is effective January 1, 2026. On December 17, 2025, the full Senate voted by a vote of 77-20 to pass the reconciled version of the National Defense Authorization Act for Fiscal Year 2026. The bill will now head to the President for final signature. The pay raise for the year 2026 will be effective January 1, 2026. On December 10, 2025, the House of representatives passed the reconciled version of the 2026 National Defense Authorization Act by a vote of 312 to 112. The full Senate is expected to also pass the bill. Once they do, it will go to President Trump for final approval. On October 9, 2025, the full Senate voted to approve their version of the National Defense Authorization Act for 2026 by a vote of

RAW → ATOMIZED

[S9] 2026 Active Duty Officer monthly basic pay for 0-8 is \$14,645 for over 4 years of service.
[S9] 2026 Active Duty Officer monthly basic pay for 0-9 is \$16,818 for 8 years of service.
[S9] 2026 Active Duty Officer monthly basic pay for 0-10 is \$18,999.
[S8] Aviation Incentive Pays include Army, Navy, and Air Force.
[S8] Monthly Air Force Aviation Incentive Pay Rates are available.
[S10] Monthly Air Force Aviation Incentive Pay Rates: 2 years of aviation service or less is \$150; over 2 years is \$250; over 6 years is \$700; over 12 years is \$1,000; over 22 or 24 years is \$450.
[S11] Aircrew members receive flight pay, also known as Aviation Career Incentive Pay (ACIP), which starts at \$150 per month and increases to \$1,000.

preview · ~1392 tokens

Done (atomize + evict raw)

Re-atomize

Cancel

Atomize

Evict raw context

Here's what happens when you click the "View" button on the context data. It shows you the raw evidence it collects from your local files, memory or web search.

- You can choose to atomize the raw data. You can see that the un-atomized raw data is estimated around ~15.8k tokens, after atomization it shrinks to ~1.4k tokens. This allows us to continue a conversation on light tokens while retaining the most important facts.
- The atomized version is actually manually editable. If you don't like how the AI used the raw data and what it extracted from it, you can re-atomize or edit manually and insert missing data or delete irrelevant things. Essentially you min max token window to get the best accuracy in as few tokens as possible.

Memory System

The screenshot displays the AI interface's memory system. On the left, a sidebar lists memory items with dates. The main area shows a chat history with a query about F1 standings and a detailed response with sources and metrics.

Memory List (Left Sidebar):

- Consolidate & Clean Memories
- Select #2
- The user expects context to the point answers to their questions. (Aug 19, 2023)
- The user desires repeat answers and enriched information. (Aug 19, 2023)
- The user's favorite color is orange. (Aug 19, 2023)
- The user's favorite F1 driver is Kimi Antonelli. (Aug 19, 2023)
- Test manual entry. (Aug 19, 2023)
- Locally can browse the web via go webserver connecting via locally browser extension on the user's browser. (Aug 19, 2023)
- The user has built a fully local AI assistant called Locally. (Aug 19, 2023)

Chat History (Main Area):

Query: "what's the f1 standings today"

Response:

A couple of notes:

- *Exact points for Hamilton, Ocon, Alonso and the others weren't shown in the retrieved standings tables, so I haven't guessed them. (For reference, Ferrari's 330 team total minus Leclerc's 155 would imply Hamilton on ~183)
- The individual driver points don't add up exactly to some team totals (e.g. Red Bull, Racing Bulls), which is consistent with mid-season driver moves — Lawson's points are split across teams, and a new Cadillac team (Bottas/Perez) is in the field this year.

Sources:

- 2023 DRIVERS STANDINGS
- 2023 F1 DRIVERS STANDINGS - ESPN
- 2023 F1 STANDINGS
- F1 constructor standings 2023 | racingnews365

Metrics:

2 calls - 121.2s - 1777 50.2s - 100m 76.2s - 28 tokens - 85 cached

System Health (Bottom Left):

- Database: Online
- Cache (Redis): Online
- OpenAI 3.5 Turbo: Online

Chat Input (Bottom):

Message AI Assistant... Send

This is the memory view. Umm, it's arguably the weakest part of the system as I haven't given it much attention and I found the best use for it for me is to ask it to remember stuff for me. Phone numbers, emails of people etc.

The screenshot displays the Gemini AI web interface. On the left is a sidebar with a 'New Chat' button and a list of previous conversations. The main chat area shows a query: 'who was it that wanted to see locally?'. Below the query, the 'Execution Trace' is visible, detailing the AI's reasoning process. This includes searching files and memories, assigning a title, and identifying the person from memories. Below the trace, a table of metrics is shown, detailing the performance of the AI's reasoning process.

Execution Trace

- ✓ Searching files & memories for "who wanted to see locally, locally"
- ✓ title assigned: "who was it that wanted to see locally?"
- ✓ context assembled with 1 prior message
- ✓ completion - searched files & memories for "who wanted to see locally, locally"

Thought Process Complete

- If a file could answer, suggest "Append to Chat" (not applicable since it's in Memories).
- Identify the person from memories: reddit user CeleryVids-4075.

Reddit user CeleryVids-4075 is interested in checking out locally.


Metrics

call	time	profile	think	text
first pass	140ms	230ms - 768 tok	95ms - 181 tok	-
answer	241ms	225ms - 187 tok	182ms - 211 tok	154ms - 17 tok

Here you can see how it works and how I personally use it. I want to and I will eventually expand on the memory system. I want to categorize and inject memories such as “facts about the model” so the model knows about itself, “style” so the model understands how I like it to respond to me, “facts about the user” so it remembers stuff specifically related to me and assists me better.


Right now the way it works is that based on your query it selects a few memory BM25 top ranking bits and pieces and injects them sparsely, which is alright but I haven’t gotten to implement it in the usual scrutinized way I implement stuff.

Emails



This is the “Mails” view. You can Add accounts here (with app password and SMTP/IMAP configs)

- The Generate Daily Briefing button essentially does what it says. It goes through each email inbox that’s hooked up and reads 5 emails at a time, sums them and then gives you an overview of what’s going on along with a button that takes you to the exact email referenced. You can also reply to an email straight from here and there is a button to make the AI draft your response for you.



Here for example I give the AI a hint in how to respond/my general attitude and then I click “AI ASSIST”, this is what happens next:

Re: [openai/codex] Auto-compaction silently discards all conversation history since 0.145.0 (Issue #36642)

FROM: Grzegorz Kaczmarek (notifications@github.com)

DATE: 2026-08-26 16:44:34

Repro 2: task disappears mid-edit

A later follow-up reported a concrete runtime assertion from an unlocked simulation-speed feature. Luna correctly diagnosed the exact problem: a value of 25 was reaching a vanilla API that asserts 1..20. It explicitly announced the intended prefix-clamp/postfix-restore fix and then modified the source accordingly.

Immediately before the next compaction:

```
input_tokens: 232179
model_context_window: 258400
```

Again, the next record was a type: "compacted" event with:

```
payload.message == ""
```

and the replacement history again lacked the current user request/task state. Four seconds later Luna replied:

Codex then marked the turn complete, despite the worktree having just been modified and no build/test validation having happened yet.

Additional detail / possible confounder

I do have the experimental token-budget feature enabled with an earlier reminder (~30k tokens remaining), plus AGENTS instructions to checkpoint task state. The reminder itself worked and, on an earlier turn, Luna created a .codex/TASK_STATE.md checkpoint. However, these two failures happened when automatic compaction itself fired, and the observable failure is specifically that the resulting compaction payload had an empty message and replacement history with no live task/history.

so the custom reminder may affect when the window approaches rollover, but it does not explain why the compaction handoff contains no task summary at all.

This looks extremely similar to the failure described here, except it is still reproducible for me on 0.150.0-alpha.0, on Windows/Desktop, and with Luna rather than Sol. The second reproduction also shows that this can leave a task half-modified and then falsely complete after the model loses its state.

—

Reply to this email directly, [view it on GitHub](#), or [unsubscribe](#).

Triage notifications, keep track of coding agent tasks and review pull requests on the go with GitHub Mobile for [iOS](#) and [Android](#). Download it today!

You are receiving this because you are subscribed to this thread.

TRANSMISSION OUTPUT

RECIPIENT HEADER

notifications@github.com

TRANSMISSION SUBJECT

Re: [openai/codex] Auto-compaction silently discards all conversation history since 0.145.0 (Issue #36642)

MESSAGE CRYPT

Thank you for providing such a detailed and rigorous reproduction of this issue.

The granularity of your report—specifically the telemetry data and the breakdown of the "compacted" event payload—is incredibly helpful. Your observation that the "replacement_history" is being stripped of live user requests, tool outputs, and task summaries, leaving only static context, points to a significant failure in the compaction handoff logic.

The impact you've described—where the model loses all context of the active task or prematurely marks tasks as complete despite pending worktree modifications—is a critical issue. Given that this is occurring in "0.150.0-alpha.0" on Windows with Luna, we will prioritize an investigation into the compaction mechanism to determine why the task summary is being omitted from the payload.

Thank you again for your thoroughness and for helping us improve the stability of Codex.


☆ AI ASSIST

Cancel


Reply

Honestly I don't actually use this feature, as I don't think people deserve AI responses but that's up to you. It's in there and it's possible and sometimes it might be justified.

Here's the email briefing:



- You get a sneak peak for each email that was considered relevant by the AI along with a button to go to the source immediately.
- It also goes through your calendar and gives you the rundown on what's up and coming.
- If any info from your email is requiring calendar entries then it suggests making those entries as well.



- It is using todoist on the back-end which you can download on your phone and works well. It was the only free option I could find so that's what it is. Make a todoist account, click settings and add your token to the .env through the UI and you're hooked up to a working calendar. It's not a google calendar and that's a good thing for me.

- In this image based on the email the AI is suggesting you create a reminder. This is really useful in my opinion and after I click that button I have it scheduled and I have reminders so it's awesome.
- There's also conflict checking. If you're trying to schedule overlapping reminders then it will tell you and allow you to create overlapping events anyway.

Settings

The screenshot shows the 'Environment Setup' dialog box with the following configuration options:

MODEL	CTX SIZE
Gemma 4 26B QAT — 160,000 ctx	160000
Gemma 4 26B QAT — 4,096 ctx	
Gemma 4 26B IQ4_NL (unsloth) — 8,192 ctx	
Gemma 4 E4B IQ4_NL (unsloth) — 65,000 ctx	
Gemma 4 E2B Q4 — 65,000 ctx	
Qwen3.5 9B IQ4_NL (unsloth) — 65,000 ctx	
12GB+	
Gemma 4 E4B Q8 — 40,960 ctx	
Gemma 4 E4B Q4 — 50,000 ctx	
Gemma 4 E2B IQ4_NL (unsloth) — 45,000 ctx	
Bonsai 27B Q1_0 — 50,000 ctx	
Gemma 4 26B IQ4_NL (unsloth) — 4,096 ctx	
Gemma 4 E4B IQ4_NL (unsloth) — 50,000 ctx	
Gemma 4 E2B Q4 — 45,000 ctx	
Qwen3.5 9B IQ4_NL (unsloth) — 32,768 ctx	
8GB+	
Gemma 4 E4B Q8 — 8,192 ctx	
Gemma 4 E4B Q4 — 12,288 ctx	
Gemma 4 E2B IQ4_NL (unsloth) — 45,000 ctx	
Gemma 4 E4B IQ4_NL (unsloth) — 16,384 ctx	
Gemma 4 E2B Q4 — 45,000 ctx	

At the bottom right, there is a 'Save Configuration' button.

You can select your model and CTX is essentially the context token size (token window).

×

 **Environment Setup**

Configure your local API connections and limits (.env)

MODEL

Gemma 4 26B QAT — 160,000 ctx

▼

CTX SIZE

160000

MAX MEMORIES LIMIT

500

DEFAULT CHAT TEMP

1

AGENT DECIDER TEMP

0.3

AGENT MEMORY SELECTOR TEMP

0.5

AGENT CONDENSER TEMP

0.5

AGENT EXTRACTOR TEMP

Cancel

Save Configuration

More general settings like max memory cap, default chat temperature etc.

×

 **Environment Setup**

Configure your local API connections and limits (.env)

MODEL

Gemma 4 26B QAT — 160,000 ctx

CTX SIZE

160000

TZ

Europe/Vienna

LOG LEVEL

info

CALENDAR FETCH LIMIT

5

MEMORY EXTRACTION THRESHOLD TOKENS

63141

TODOIST API KEY

Cancel

Save Configuration


This is where you put your todoist key to connect to your calendar.

Jobs

This doesn't mean chronjobs or agentic tasks queued. This is quite literally jobs.

- You can enter search urls into the registry manually and the app scrapes the job websites for x amount of job titles, decides if your uploaded resume is a match for the job and cross checks your profile for data about your preferences.
- I built this because I'm looking for a job and a lot of my friends are also looking for one. This works for any job, any resume etc.
- You can see in the logs why something was selected, why something was discarded or if it couldn't parse something.

- You can enter search urls into the registry manually and the app scrapes the job websites for x amount of job titles, decides if your uploaded resume is a match for the job and cross checks your profile for data about your preferences.
- I built this because I'm looking for a job and a lot of my friends are also looking for one. This works for any job, any resume etc.
- You can see in the logs why something was selected, why something was discarded or if it couldn't parse something.



Here you can see the profile settings, these are your preferences essentially.

That's about it. I hope you can find some cool ideas in here or if you wanna fork it and work on it for yourself that also works.

Features I plan to add in the future:

- Watch YouTube or other video links for me. You paste a link, it downloads the audio of the video, runs transcription with whisper running on cpu and sums up the transcription to answer your questions about some video you don't have time to watch but are interested in based on the title or thumbnail.
- Link context. You paste a link to a reddit post, blog, some website whatever it is and it parses it and can talk to you knowing exactly what's on the page you linked.
- Categorization for memories. Memories about itself, memories about the user, memories about the style that it should use when interacting with the user and so on.
- Auto-draft applications to jobs via the email view. This means it finds a job, finds the recruiters information, drafts an email and notifies you that the email is ready to be sent, you just click a send button.

That's about it. If you have questions or want more in-depth implementation details feel free to ask. 😊